
IM0802 – Responsible Artificial Intelligence

Assignment 1: Deepfakes in time of conflict or war

Harman Preet Singh¹
852863316

1 Introduction

A United Nations Security Council (UNSC) meeting is underway to vote on a ceasefire agreement for a conflict that has been going on for years. Both sides are accusing the other party of war crimes without concrete. Online, a video surfaces showing soldiers from one party carrying out organised killings of unarmed civilians. The video quickly goes viral, and is shared by news outlets across social media. Diplomats condemn the killings and the UNSC imposes economic sanctions. The accused nation denies the allegations and claims that the video is fabricated. However, the damage is already done. Diplomatic relations have been lost and strict economic sanctions are in place. The latter eventually leads to the accused nation's economy collapsing and a forced withdrawal from the conflict.

However, there is more to the story. The killings never happened. The video was a deepfake, generated by the government of the opposing nation to discredit the accused nation. Generative artificial intelligence (Gen AI) has made significant advancements, with recently footage generated by Gen AI being almost indistinguishable from reality. Until not very long ago, deepfakes were not very convincing, and were mostly used for entertainment purposes. But in 2026, the advancements of generative artificial intelligence could very well be used to fabricate a video with deepfake technology to influence international diplomacy. This essay will explore the ethical implications of deepfakes in time of conflict or war.

2 Issue

Gen AI is the primary tool used in deepfake technologies, to manipulate images, videos or audio to convincingly replace one person's likeness with another's [1]. This is, besides for entertainment purposes, often done for malicious purposes. Early deepfakes were easy to identify because they were not very realistic. However, advances in Gen AI have changed that. Now, deepfakes are harder to detect and that makes them a powerful tool for disinformation campaigns, for example in conflict zones [1, 2].

Disinformation using Gen AI is not a theoretical concern. It has been actively deployed in real-world conflict zones, like Ukraine. The DFRLab report [3] documents multiple instances where deepfakes were used to spread false information. Examples include fabricated videos of political leaders making inflammatory statements or staged military actions. In the case of Ukraine, deepfakes have been used to create confusion, with one video¹ showing a fabricated speech by Ukrainian president, Volodymyr Zelenskyy, calling his own troops to surrender themselves.

The consequences of deepfakes in conflict zones do not just affect the government and soldiers. Civilian populations can be misled by fabricated videos, which often leads to panic or misguided actions. Opposing governments may use deepfakes as propaganda tools to justify aggressive policies. International bodies like the UN may struggle to respond effectively when visual evidence is no longer trustworthy. Journalists face challenges in verifying sources and maintaining credibility. The

¹<https://www.youtube.com/watch?v=X17yrEV5sl4>

general public may become increasingly skeptical of all media because trust in information sources has eroded [1, 2].

Wartime propaganda and deception are not new phenomena. They have been part of human conflict for centuries. However, AI changes the scale, accessibility, and believability of disinformation. Deepfakes were easily discernible in the past, but now increasingly people are struggling to distinguish real from fake. This creates a never-before-seen ethical challenge because it undermines trust in visual evidence. As Kenderdine and Hibberd [4] argue, deepfakes erode our cultural and epistemic trust in images and videos, which have traditionally been powerful tools for conveying truth.

Pre-existing societal conditions such as political polarization, weakened media literacy, and fractured international trust amplify the impact of deepfakes. In highly polarised societies, people may be more likely to accept deepfakes that align with their beliefs and reject those that do not. Countries actively in conflict are inherently polarised, making them susceptible to deepfakes that reinforce existing narratives. Weakened media literacy means that many individuals lack the skills to critically evaluate sources. This makes them more susceptible to manipulation. Additionally, fractured international trust can hinder coordinated responses to disinformation campaigns.

3 Solution

There are various technical approaches to detect deepfakes. Forensic analysis involves examining the digital artifacts left by deepfake generation processes, such as inconsistencies in lighting, shadows, or facial movements. Authentication watermarking involves embedding a unique identifier into original content that can be verified later to confirm authenticity. However, these methods have limitations. As Mah [1] points out, deep learning and NLP techniques used for detection often lag behind the advancements in deepfake generation, making it a constant cat-and-mouse game between creators and detectors.

Ethical frameworks for responsible AI development, such as those outlined by Akkuş [2], emphasize principles like accountability, transparency, and fairness. In the context of deepfakes used in conflict, these principles raise critical questions about who is responsible for the harm caused. If a deepfake video incites violence or spreads misinformation, should the creators of the deepfake be held accountable? What about platforms that host such content, or do not take sufficient measures to fight misinformation? Sole responsibility is difficult to assign in the digital age, where content can be easily shared and manipulated.

Certain regulatory responses have been proposed to address the challenges posed by deepfakes. The EU AI Act [5], for example, aims to establish a legal framework for AI applications, including provisions for high-risk AI systems like deepfake generators. However, as Akkuş [2] argues, existing legal frameworks may not be sufficient to address the unique challenges of deepfakes in conflict contexts. International Humanitarian Law (IHL) may not adequately cover the use of deepfakes as a weapon of war, and there is a need for new regulations that specifically address the ethical and legal implications of AI-generated disinformation in conflict zones.

Beyond technical and regulatory solutions, media literacy programs can be effective to educate the general public about the existence and risks of deepfakes. NGOs and journalism standards organizations, such as the Content Authenticity Initiative, work to establish best practices for verifying content and maintaining journalistic integrity. International standards bodies like IEEE can also play a role in developing guidelines for responsible AI use. A collaboration between the technical, regulatory, and societal mechanisms is essential to dampen the amplifying effect of societal factors like political polarization and weakened media literacy that make populations more vulnerable to deepfake disinformation.

It's my opinion that deepfakes are in fact an issue in that not everyone has the technical expertise to identify them. While I myself, and those I spend time with are mostly aware of deepfakes and their risks, I worry about the general public's ability in this new landscape of disinformation. With how polarised the world is, combined with the fact that every social media platform has a tailored algorithm that feeds us content we are likely to blindly agree with, I think deepfakes are a real threat to the political climate. I don't think any single solution is sufficient to address this issue. In my opinion, and that of many others, a multi-layered approach that combines technical detection methods, regulatory frameworks, and societal mechanisms is necessary.

4 Conclusions

The fictional ethical dilemma from the introduction section is not far-fetched given the documented cases discussed in the Issue section. Generative Artificial Intelligence has transformed the old problem of wartime propaganda into something both qualitatively and quantitatively different. The erosion of trust in visual trust is a fundamentally new issue, even if the underlying problem of propaganda is not. However, this issue is not caused by AI alone. It merely amplifies societal conditions like polarisation or a weak media literacy.

According to Mah [1] Detection tools exist but they are reactive by nature, meaning they lag behind the development of increasingly improving deepfake tools. From a legal perspective, Akkuş [2] reveals a significant gap in International Humanitarian Law, which was not designed with AI-generated deception in mind. Existing frameworks, like the Geneva Convention, are difficult to apply in the case of deepfakes. The DFRLab reports that even well-documented cases of deepfake-based disinformation in Ukraine did not prevent confusion among the public [3]. This shows that the effort of current societal mechanisms, like media literacy or fact-checking organisations, are not sufficient and is being outpaced by the speed at which deepfakes spread over social media.

I don't believe a total change in approach is necessary, but rather that the existing solutions need to be reviewed and adjusted accordingly. Technical detection needs to be embedded into platforms at the point of distribution, and progress is being made in this direction. International Humanitarian Law should be updated to explicitly address the issue of AI-generated deception. Finally, more effort should be invested in increasing media literacy, especially in vulnerable regions such as conflict zones.

The pace of AI development will only intensify the issue of deepfakes, so coordinated action is urgently needed rather than optional.

References

- [1] Pascal Muam Mah. The role of deepfake, deception, and disinformation in conflict zones based on dl for nlp: A critical ai-era perspective. *Comunicar*, 34(84):228–246, 2026. doi: 10.5281/zenodo.18115680. URL <https://doi.org/10.5281/zenodo.18115680>.
- [2] Berkant Akkuş. Deepfakes and the geneva conventions: Does deceptive ai-generated misinformation directed at an enemy during armed conflict violate international humanitarian law? a critical discussion. *Laws*, 14(6), 2025. ISSN 2075-471X. doi: 10.3390/laws14060083. URL <https://www.mdpi.com/2075-471X/14/6/83>.
- [3] Roman Osadchuk. Ai tools usage for disinformation in the war in ukraine, Jul 9 2024. URL <https://dfrlab.org/2024/07/09/ai-tools-usage-for-disinformation-in-the-war-in-ukraine/>. Digital Forensic Research Lab (DFRLab) report.
- [4] Sarah Kenderdine and Lily Hibberd. *Deep Fakes: A Critical Lexicon for Digital Museology*. Routledge, 10 2025. ISBN 9781003334453. doi: 10.4324/9781003334453.
- [5] European Parliament and Council of the European Union. Regulation (eu) 2024/1689 of the european parliament and of the council of 13 june 2024 laying down harmonised rules on artificial intelligence (artificial intelligence act). <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024R1689>, 2024. Official Journal of the European Union, OJ L 2024/1689, 12 July 2024.