
IM0802 – Responsible Artificial Intelligence

Assignment 2: “Paper Title Goes Here”

Alexander Soerjo
852767169

Harman Singh
852863316

Abstract

Your abstract. Your abstract.

1 Introduction 350

Israel, which allocates 8.8% of its GDP to defense [1], invests heavily in the development of autonomous weapon systems (AWS), with their Prime Minister announcing a \$100 billion investment in homegrown arms production last December [2]. Israel’s most well-known AWS is the Iron Dome air defense system, which intercepts incoming projectiles in real time. Building on this foundation in artificial intelligence (AI)-assisted defense technology, Israel has expanded into offensive targeting applications such as “Habsora”.

Habsora (“The Gospel”) is an Israeli Defense Forces (IDF) AI-targeting system developed primarily by Unit 8200 within the Military Intelligence Directorate. First deployed during the May 2021 Operation “Guardian of the Walls”, it generates target recommendations using a multimodal intelligence pipeline that encompasses signal intelligence, telephone metadata, satellite and drone imagery, social graph analysis, and life-pattern modelling [3, 4]. The system is not a single model, but an ensemble of multiple specialised algorithms fused together [4]. During the 2023-2024 Gaza conflict, Habsora was reported to generate over one hundred target recommendations per day, which in prior years required approximately twelve months of human analysts [3, 5].

Habsora’s AI capability is the probabilistic classification and prioritisation of locations and individuals as military objectives. It assigns confidence scores derived from large-scale pattern matching across a shared intelligence pool, rather than positive identification in the traditional intelligence analysis [6]. Military commands nominally authorise each recommendation, yet their review window documented for each target was approximately twenty seconds [3].

This creates a stakeholder dynamic with both ethical and legal significance. It’s output influencing both IDF commands, who operate under high cognitive load and time pressure; Palestinian civilians in the target environment bearing the lethal consequences and the broader international community which depends on International Humanitarian Law (IHL) compliance for the legitimacy of armed conflict. Steen et al. [7] offer a framework for this tension. They propose a two-axis model of Meaningful Human Control (MHC) which distinguishes between high automation and genuine human control, which both require adequate information and real freedom of action. The central argument of this report is that Habsora structurally undermines the second condition, and that this gap carries cascading implications across the technical, sociotechnical, and governance layers identified by Verdiesen et al. [8] in their Comprehensive Human Oversight Framework (CHOF).

2 Ethical Issue Identification 750

	Before	During	After
Technical	Partial Built; no XAI embedded	Absent No reasoning chain shown	Absent No performance audit
Sociotechnical	Partial No bias mitigation training	Absent 20-sec rubber stamp dynamic	Absent No incident review process
Governance	Absent No EIA or legal pre-approval	Absent No IHL compliance check	Absent No audit or accountability

Figure 1: The CHOF framework for structuring ethical analysis across governance, socio-technical, and technical layers. Two yellow cells indicate where Habsora can be given partial compliance. The red cells indicate absence.

Verdiesen et al. [8]’s Comprehensive Human Oversight Framework (CHOF) was developed to assess accountability for autonomous systems whose internals are not directly inspectable. It divides oversight into a 3×3 matrix, before, during and after deployment across a technical, socio-technical and governance mapping. The resulting matrix surfaces gaps which remain hidden when ‘human oversight’ would be treated as a single requirement. Applying this framework to Habsora’s implementation we get Figure 1.

2.1 Technical Layer

Habsora has been described as an ensemble of deep learning models processing multimodal data [3, 4]. Such architectures are inherently opaque, Arrieta et al. [9] characterise this as the accuracy/interpretability trade-off, where deep ensembles maximize predictive performance at the cost of transparency. While post-hoc explainability methods exist that could be applied to such systems, reports on Habsora’s use indicate that no such methods are applied as the operators point of decision.

This opacity directly links to the legal consequences. Schmitt and Thurnher [6] establish that IHL’s distinction principle requires positive identification of a target as being a military objective. In conjunction with the proportionality assessment being understood by the user. Without any interpretability the legal review is effectively empty.

The above failures map directly into the highest-ranked principles identified by Khan et al. [10]. Both **P-01** transparency and **P-03** accountability, further compounded by their challenges: lack of technical understanding among decision-makers **CH-06** and the absence of binding legal frameworks **CH09**. While this might not be a failure of intent, it is a structural property of opaque architectures. Combined with lacking operator-facing interfaces, this creates a system which can not provide the context for understanding.

2.2 Socio-technical

The Value Sensitive Design (VSD) framework addresses indirect stakeholders, those affected without directly using the system, such as civilians in the target environment. Palestinian civilians were not involved in its design, and their values were not elicited to ground the system’s targeting logic. They remain non-consenting parties who face the consequences of the system’s use without recourse. There is a large asymmetry in consequences between those operating the system and those subjected to it.

For IDF commanders, the conditions necessary for genuine human oversight are structurally not present. Steen et al. [7] argue that meaningful human control requires both sufficient information and meaningful freedom of action. Habsora undermines both: commanders receive recommendations

without transparency into the underlying reasoning, and the short decision windows of the system leave almost no room for critical evaluation. Miller [11] shows that from social science, without understanding what went into a decision, the explainee cannot meaningfully interrogate it. Combined with the large amount of targets the system generates, this reduces what would normally require extended deliberation and proportionality assessments to a button press. Former IDF commanders have described this process as ‘rubber-stamping’ [12]. This confirms the presence of automation bias, which is when operators rely too much on AI outputs under time pressure [13].

The consequence is that IDF commanders formally bear moral and legal responsibility for strikes they cannot meaningfully evaluate, while the system’s opacity diffuses accountability in practice. Targeting decisions that once demanded careful human judgement are now effectively delegated to an algorithm. This prevents the human control which the IHL demands, and that Steen et al. [7] identify as a precondition for attributing responsibility.

2.3 Governance

Article 36 of Protocol I to the Geneva Conventions requires countries to perform legal reviews of new weapons systems to determine whether their use would be prohibited under international law [14]. Israel is one of the few countries that have not ratified Protocol I [15]. UNESCO maintains that delegating life-or-death decisions to an AI system is itself forbidden [16]. Yet, no accountability framework exists to enforce this in the case of Habsora **TODO EVIDENCE NEEDED**.

Once deployed, establishing accountability becomes difficult. Under IHL, commanders are legally responsible for how weapons are used, even when autonomous functions are involved [17]. They are judged on what they knew at the time of the decision [18]. In practice, responsibility is spread across developers, the algorithm, the receiving IDF officer, and the approving commander [16]. The Lavender system, the human-targeting extension of Habsora, has a documented 10% error tolerance in targeting recommendations. XYZ This illustrates the problem: this margin makes civilian casualties inevitable. In the midst of war, real-time monitoring is not easy to pull off, which compounds the issue.

After deployment, we find that review processes are largely absent. The system’s opacity makes it difficult to determine whether casualties resulted from an AI recommendation or a human override. Palestinian civilians, who the most directly affected stakeholders, were never consulted. No structured ethical assessment, such as ‘value sensitive design’ (VSD), was conducted.

3 Responsible Solution Design 900

De literatuur met 'NEW' erbij staan reeds in de bronnen van de case (in het bestand RAI Assignment 2 Description) maar hebben we nog niet besproken omdat we nog niet bezig waren met solutions

3.1 Technical layer (XAI, interpretability)

- **Arrieta et al. [9]**: post-hoc solutions (SHAP, LIME, attention visualisation) applied to ensemble models
- **Miller [11]**: XAI
- **(NEW) Liao, Gruen & Miller (2020) [19]**: XAI question bank (Why this target? What evidence supports this recommendation?) → design requirements
- **(NEW) Morley et al. (2020) [20]**: maps ethics tools onto ML pipeline stages (data, model, deployment) so we can use this to say where and how specific XAI tools can be added before/during/after model development

3.2 Socio-technical layer (MHC, automation bias, VSD, stakeholder inclusion)

- **Steen et al. [7]**: The 2 MHC conditions (sufficient information, meaningful freedom of action) → design requirements
- **Parasuraman and Manzey [13]**: VSD's 'tripartite' methodology (conceptual, empirical, technical) → concrete process for how stakeholders should have been consulted (Palestinians were not consulted)
- **(NEW) Amershi et al. (2019) [21]**: human-AI interaction guidelines (calibration of trust, communicating uncertainty, supporting error recovery) → design principles

3.3 Governance (accountability, legal frameworks, organisational response)

- **Verdiesen et al. [8]**: addresses the empty cell in Figure 1 → required governance interventions (pre-deployment EIA, real-time IHL compliance, post-deployment audit)
- **(NEW) Stahl et al. (2022) [22]**: Use the 3-level mitigation (policy-level strategies, guidance mechanisms, corporate governance) to categorise governance solutions
- **Wasilow and Thorpe [16]**: Canada has pre-deployment legal assessment protocols (contrast with Israel that does not, because it does not follow Article 36 review process)
- **Khan et al. [10]**: P-01/P-03/CH-06/CH-09 → we can use this to motivate specific governance solutions

References

- [1] Stockholm International Peace Research Institute. Trends in world military expenditure, 2024. SIPRI, April 2025. URL https://www.sipri.org/sites/default/files/2025-04/2504_fs_millex_2024.pdf. SIPRI Fact Sheet. Accessed: 07 April 2026.
- [2] Times of Israel Staff. Israel investing more than \$100 billion into independent arms industry, Netanyahu reveals. *The Times of Israel*, December 2025. URL <https://www.timesofisrael.com/israel-investing-more-than-100-billion-into-independent-arms-industry-netanyahu-reveals/>. Accessed: 12 April 2026.
- [3] Yuval Abraham. “lavender”: the ai machine directing israel’s bombing spree in gaza, Apr 2024. URL <https://www.972mag.com/lavender-ai-israeli-army-gaza/>.
- [4] Elizabeth Dwoskin. Israel built an “ai factory” for war. it unleashed it in gaza., Dec 2024. URL <https://www.washingtonpost.com/technology/2024/12/29/ai-israel-war-gaza-idf/>.
- [5] Ynet. “man replaced by machine”: Is the use of ai undermining the idf’s intelligence capabilities?, Dec 2024. URL <https://www.ynetnews.com/article/sjyhbni1kx>.
- [6] Michael N. Schmitt and Jeffrey S. Thurnher. “out of the loop”: Autonomous weapon systems and the law of armed conflict. *Harvard National Security Journal*, 4(2):231–281, 2013. URL <https://centaur.reading.ac.uk/89863/>.
- [7] Marc Steen, Jurriaan Diggelen, Tjerk Timan, and Nanda Stap. Meaningful human control of drones: exploring human–machine teaming, informed by four different ethical perspectives. *AI and Ethics*, 3, 05 2022. doi: 10.1007/s43681-022-00168-2.
- [8] Ilse Verdiesen, Andrea Aler Tubella, and Virginia Dignum. Integrating comprehensive human oversight in drone deployment: A conceptual framework applied to the case of military surveillance drones. *Information*, 12(9), 2021. ISSN 2078-2489. doi: 10.3390/info12090385. URL <https://www.mdpi.com/2078-2489/12/9/385>.
- [9] Alejandro Barredo Arrieta, Natalia Díaz-Rodríguez, Javier Del Ser, Adrien Bannetot, Siham Tabik, Alberto Barbado, Salvador García, Sergio Gil-López, Daniel Molina, Richard Benjamins, Raja Chatila, and Francisco Herrera. Explainable artificial intelligence (xai): Concepts, taxonomies, opportunities and challenges toward responsible ai, 2019. URL <https://arxiv.org/abs/1910.10045>.
- [10] Arif Ali Khan, Sher Badshah, Peng Liang, Bilal Khan, Muhammad Waseem, Mahmood Niazi, and Muhammad Azeem Akbar. Ethics of ai: A systematic literature review of principles and challenges, 2021. URL <https://arxiv.org/abs/2109.07906>.
- [11] Tim Miller. Explanation in artificial intelligence: Insights from the social sciences, 2018. URL <https://arxiv.org/abs/1706.07269>.
- [12] Atay Kozlovski. When algorithms decide who is a target: Idf’s use of ai in gaza. *Tech Policy Press*, May 2024. URL <https://www.techpolicy.press/when-algorithms-decide-who-is-a-target-idfs-use-of-ai-in-gaza/>.
- [13] Raja Parasuraman and Dietrich H. Manzey. Complacency and bias in human use of automation: An attentional integration. *Human Factors: The Journal of the Human Factors and Ergonomics Society*, 52(3):381–410, 2010. doi: 10.1177/0018720810376055.
- [14] International Committee of the Red Cross. Protocol additional to the geneva conventions of 12 august 1949, and relating to the protection of victims of international armed conflicts (protocol i), 8 june 1977, article 36 - new weapons. <https://ihl-databases.icrc.org/en/ihl-treaties/api-1977/article-36>, 1977. ICRC Database, Treaties, States Parties and Commentaries. Last accessed: 22 April 2026.
- [15] International Committee of the Red Cross. State parties to Additional Protocol I of 1977. IHL Treaties Database, 2024. URL <https://ihl-databases.icrc.org/en/ihl-treaties/api-1977/state-parties>.
- [16] Sherry Wasilow and Joelle B. Thorpe. Artificial intelligence, robotics, ethics, and the military: A canadian perspective. *AI Magazine*, 40(1):37–48, 2019. doi: <https://doi.org/10.1609/aimag.v40i1.2848>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1609/aimag.v40i1.2848>.

- [17] International Committee of the Red Cross. Protocol I additional to the Geneva Conventions of 12 august 1949, article 86: Failure to act. IHL Treaties Database, 1977. URL <https://ihl-databases.icrc.org/en/ihl-treaties/api-1977/article-86>.
- [18] Jean-Marie Henckaerts and Louise Doswald-Beck. *Customary International Humanitarian Law*, volume 1. Cambridge University Press, Cambridge, 2005.
- [19] Q. Vera Liao, Daniel Gruen, and Sarah Miller. Questioning the AI: Informing design practices for explainable AI user experiences. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*, pages 1–15, New York, NY, USA, 2020. ACM. doi: 10.1145/3313831.3376590.
- [20] Jessica Morley, Luciano Floridi, Libby Kinsey, and Anat Elhalal. From what to how: An initial review of publicly available AI ethics tools, methods and research to translate principles into practices. *Science and Engineering Ethics*, 26(4):2141–2168, 2020. doi: 10.1007/s11948-019-00165-5.
- [21] Saleema Amershi, Dan Weld, Mihaela Vorvoreanu, Adam Fourney, Besmira Nushi, Penny Collisson, Jina Suh, Shamsi Iqbal, Paul N. Bennett, Kori Inkpen, Jaime Teevan, Ruth Kikin-Gil, and Eric Horvitz. Guidelines for human-AI interaction. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, pages 1–13, New York, NY, USA, 2019. ACM. doi: 10.1145/3290605.3300233.
- [22] Bernd Carsten Stahl, Josephina Antoniou, Mark Ryan, Kevin Macnish, and Tilimbe Jiya. Organisational responses to the ethical issues of artificial intelligence. *AI & Society*, 37(1): 23–37, 2022. doi: 10.1007/s00146-021-01148-6.
- [23] Batya Friedman, Peter H. Kahn, Alan Borning, and Alina Huldtgren. Value sensitive design and information systems. In Neelke Doorn, Daan Schuurbiers, Ibo van de Poel, and Michael E. Gorman, editors, *Early Engagement and New Technologies: Opening Up the Laboratory*, pages 55–95. Springer, Dordrecht, 2013. doi: 10.1007/978-94-007-7844-3_4.
- [24] Batya Friedman, Jr. Kahn, Peter H., and Alan Borning. Value sensitive design and information systems. In Kenneth E. Himma and Herman T. Tavani, editors, *The Handbook of Information and Computer Ethics*, pages 69–101. John Wiley & Sons, Hoboken, NJ, 2008. doi: 10.1002/9780470281819.ch4.

A Appendix

Table of Abbreviations

AI	Artificial Intelligence
AWS	Autonomous Weapon Systems
CH-06	Challenge 06: Lack of technical understanding among decision-makers [10]
CH-09	Challenge 09: Absence of binding legal frameworks [10]
CHOF	Comprehensive Human Oversight Framework [8]
EIA	Ethical Impact Assessment
HCI	Human-Computer Interaction
IDF	Israel Defense Forces
IHL	International Humanitarian Law
LAWS	Lethal Autonomous Weapon Systems
LIME	Local Interpretable Model-agnostic Explanations
MHC	Meaningful Human Control [7]
ML	Machine Learning
P-01	Principle 01: Transparency [10]
P-03	Principle 03: Accountability [10]
SHAP	SHapley Additive exPlanations
UX	User Experience
VSD	Value Sensitive Design [23, 24]
XAI	Explainable Artificial Intelligence

Glossary of Terms

Term	Definition
Term	Definition in the format of a sentence
Term	Definition in the format of a sentence
Term	Definition in the format of a sentence